



Red Hat Reference Architecture Series

OpenShift on OpenStack

Mark Lamourine

Version 0.1, 2016-01-20




Table of Contents

Comments and Feedback	2
<i>Staying In Touch</i>	2
<i>Like us on Facebook</i>	2
<i>Follow us on Twitter</i>	2
<i>Plus us on Google+</i>	2
Introduction	3
Goals	3
OpenShift and OpenStack	3
Installation/Orchestration: <i>Heat</i>	4
Container Networking: <i>Neutron</i> and <i>Flannel</i>	4
Block Storage: <i>Cinder</i>	4
Authentication and Access Control: <i>Keystone</i>	4
Integration and Validation	4
Architecture	5
OpenStack	5
OpenShift	5
Interactions	5
Networking	5
Storage	5
Authentication	5
Preparation	6
Installation	7
References	8
Visibility and Validation	9
Appendix	10

100 East Davie Street
Raleigh NC 27601 USA
Phone: +1 919 754 3700
Phone: 888 733 4281
PO Box 13588
Research Triangle Park NC 27709 USA

Linux is a registered trademark of Linus Torvalds. Red Hat, Red Hat Enterprise Linux and the Red Hat "Shadowman" logo are registered trademarks of Red Hat, Inc. in the United States and other countries.

Ceph is a registered trademark of Red Hat, Inc.

UNIX is a registered trademark of The Open Group.

Docker and the Docker logo are trademarks or registered trademarks of Docker, Inc.

Intel, the Intel logo and Xeon are registered trademarks of Intel Corporation or its subsidiaries in the United States and other countries. All other trademarks referenced herein are the property of their respective owners.

© 2015 by Red Hat, Inc. This material may be distributed only subject to the terms and conditions set forth in the Open Publication License, V1.0 or later (the latest version is presently available at <http://www.opencontent.org/openpub/>).

The information contained herein is subject to change without notice. Red Hat, Inc. shall not be liable for technical or editorial errors or omissions contained herein.

Distribution of modified versions of this document is prohibited without the explicit permission of Red Hat Inc.

Distribution of this work or derivative of this work in any standard (paper) book form for commercial purposes is prohibited unless prior permission is obtained from Red Hat Inc.

The GPG fingerprint of the security@redhat.com key is: CA 20 86 86 2B D6 9D FC 65 F6 EC C4 21 91 80 CD DB 42 A6 0E

Send feedback to refarch-feedback@redhat.com

Comments and Feedback

In the spirit of open source, we invite anyone to provide feedback and comments on any reference architecture. Although we review our papers internally, sometimes issues or typographical errors are encountered. Feedback allows us to not only improve the quality of the papers we produce, but allows the reader to provide their thoughts on potential improvements and topic expansion to the papers. Feedback on the papers can be provided by emailing refarch-feedback@redhat.com. Please refer to the title within the email.

Staying In Touch

Join us on some of the popular social media sites where we keep our audience informed on new reference architectures as well as offer related information on things we find interesting.

Like us on Facebook

<https://www.facebook.com/rhrefarch>

Follow us on Twitter

<https://twitter.com/RedHatRefArch>

Plus us on Google+

<https://plus.google.com/u/0/b/114152126783830728030/>

Introduction

Goals

The goal of this document is to demonstrate the installation of OpenShift in an OpenStack environment with highlights on those places where OpenShift can make special use of OpenStack resources.

Ordinarily a service running on OpenStack wouldn't be aware of that fact and shouldn't need to be. OpenShift requires and provides several services that could work in concert with OpenStack to improve the quality of the user experience for operators, application developers and for application users.

This document will concentrate on using those OpenStack resources to create and run OpenShift. It will become clear that most of the actual installation process will be handled by orchestration systems, *Heat* and *Ansible*. The challenge is in choosing the correct parameters and initial environment so that the result matches expectations. Where it is useful we will show the commands which are run to execute the installation beneath the surface. We will also show how to determine that the integrated service is working as expected.

OpenShift and OpenStack

[OpenShift](#) is a premier open source *Platform as a Service* (PaaS) system. It offers developers a consistent self-service resource for containerized applications. For production services it is a stable application publication system. It has characteristics that make *High Availability* (HA) and scaling of applications under load relatively automatic.

OpenShift offers enterprise quality control and management of containerized services.

[OpenStack](#) is the leading private cloud *Infrastructure as a Service* (IaaS) system. It puts computational, networking and storage resources under software control. Using *virtual machines* (VM) as the unit of computation, *Software Defined Networks* (SDN) to connect the VMs together and with the outside world and networked filesystems such as Ceph and Gluster, OpenStack overlays a set of abstractions that allow the creation of large-scale highly flexible computer environments.

In most cases OpenShift uses the resources OpenStack provides in a totally transparent way. OpenShift is installed on VMs within the OpenStack service just as if they were on hardware. The components communicate over the network and write to storage just as if it were hardware as well. When this holds, OpenShift is entirely unaware that it is running on OpenStack. This is the purpose of OpenStack.

There are several cases though, where OpenShift offers services to its own users, OpenShift can make direct use of OpenStack resources.


Installation/Orchestration: *Heat*

Container Networking: *Neutron and Flannel*

Block Storage: *Cinder*

Authentication and Access Control: *Keystone*

Integration and Validation



Architecture

OpenStack


OpenShift

Interactions

Networking

Storage

Authentication



Preparation




Installation



References

- [OpenShift](#)
- [OpenStack](#)
 - [OpenStack Heat](#) - OpenStack Orchestration
 - [OpenStack Cinder](#) - Block Storage Service
 - [Keystone](#) - Identity Management
- [Ansible](#) - Host Orchestration



Visibility and Validation

Once Heat has done it's work and created the OpenShift service a user is going to want to verify that the installation has completed successfully and that the service is working as expected.



Appendix



